

PNsAV

Probabilistic Neuro-symbolic Argument Validation Model

Cosma Rareş-Gabriel¹, Lăzărescu Larisa-Ioana¹

¹Colegiul Național "Octavian Goga", Sibiu, România

Abstract. PNsAV is a multi-layer neuro-symbolic reasoning system designed to perform validation of argument structures extracted from natural language. A constrained Large Language Model orchestration is employed as a semantic parser, identifying argumentative units, inferential relations, and domain concepts. These structured arguments are processed by a set of formal validation layers, used for checking the structural integrity as well as the presence of cycles and graph density. We integrate a logic engine that takes into account each argument’s internal structure through Spaans’s intrinsic strength, which is determined by numeric strengths assigned to premises and rules manually according to a specified rubric, and the global relationships between arguments, to compute a converging final strength for every argument.

1 Introduction

In recent years, the rapid advancement of Large Language Models, and their deployment in reasoning environments outpaced our ability to formally validate their outputs. While the output capacity of these models is substantial, in critical fields such as law or medicine, LLMs can generate factually inconsistent content, a phenomenon termed "hallucination" (*I*), without a traceable reasoning path, thereby challenging the safety of autonomous decision-making.

Formal logic presents multiple frameworks and tools that aid proving factual information. The current methods of mapping arguments and debates lack the scalability of a neural system, as most of them are to be manually populated with data.

Current literature explores the bridge between the two, as shown in Alfano’s 2026 FAKB pipeline (15), or Pan’s 2023 Logic-LM (12), but they often fail to target the internal structure of the arguments, remaining focused on the global acceptability, being divergent in most cases from the actual strength of the arguments themselves.

We propose PN_sAV, a multi-layer end-to-end neuro-symbolic reasoning system that couples a constrained Large Language Model orchestration for natural language parsing, reducing the risk of hallucinations, with a symbolic layer of validation and logical pipeline capable of assessing both local and global structure in order to derive competent strength values starting from a manual premise and rule assessment.

2 Related Work

The development of a complete neuro-symbolic architecture capable of extracting and validating argumentative structures requires extensive research in formal logic and in natural language processing and AI. The suite of metrics for scoring reasoning proposed by Olga G. et al. in 2023 (1) proves that LLMs can exhibit specific behaviors such as incoherence, redundancy, or hallucinations. However, as identified by Miller, Nate (2026) (2), a multi-agent orchestration can be an efficient and cost-effective argument mining pipeline.

In order to categorize, analyze and express arguments mathematically, we use specific frameworks. First proposed by P. M. Dung in 1995 (3), the Abstract Argumentation Framework established the theoretical foundation for many others, such as the ASPIC⁺ framework for structured argumentation (4). The Abstract Argumentation Framework is structured on the relations between the arguments, represented as abstract nodes in a graph. Utilizing this structure, a system can compute the extension, a set of arguments which are logically justified, under a variety of formal defense semantics. ASPIC⁺ is a framework designed to augment Dung’s AAFs, offering information on the nature of the attacks and defeats as well as on the structures of the arguments, to improve the modeling of reasoning problems.

Beyond the qualitative in/out/undecided status of an argument, weighted bipolar argumentation graphs (16) assign each argument a graded strength in $[0, 1]$, propagated through attack and support relations, allowing for more fine-grained reasoning about argumentative disputes. However, these

frameworks typically leave an argument’s initial strength as an exogenous, externally supplied value. Spaans (20) addresses this gap by studying principled methods for deriving an argument’s intrinsic strength directly from the weights of the premises and rules used to construct it, rather than treating it as atomic. Separately, convergence of strength propagation on cyclic graphs remained poorly understood until Mossakowski and Neuhaus (17) introduced a modular decomposition of aggregation and influence functions and gave the first convergence and divergence results for cyclic graphs, later simplified and generalized by Potyka (19) via the contraction mapping principle, at the cost of a trade-off with open-mindedness.

Merging those two methodologies, the neural argument mining and the symbolic formal logic computations, emerges an end-to-end neuro-symbolic system that maximizes the scalability and validity of both methods. Pre-LLM neuro-symbolic systems showcase scalability issues arising from the structural dependency on problem-specific modules (5–10). Newer neuro-symbolic models, that solve the natural language parsing using LLMs, such as Logic-LM (12) and SymbCoT (11), still either lack defeasible logic (12), or use an LLM only architecture (11). Latest literature, like the FAKB pipeline by Alfano et al. in 2026 (15), presents structural flaws, more specifically in the way that they denote an argument’s strength. In the cited paper, the strength is derived from the parser LLM’s log probabilities. That causes a structural misalignment because the confidence score outputted by the LLM is based on the extraction precision rather than the argument’s actual potential value.

While alternative neuro-symbolic integration strategies are actively being explored to stabilize generative AI outputs, PNsAV uniquely focuses on an end-to-end pipeline that maps unstructured dialogue directly into weighted argumentative structures that are compositionally analyzed to provide intrinsic strength values that can be used as a basis for evaluating argumentative disputes.

3 Theoretical Basis

3.1 Dung’s Abstract Argumentation Framework (3)

Let $\mathcal{AF} = \langle \mathcal{A}, at \rangle$ be a pair where $\mathcal{A}$ is a set of arguments and at a binary relation on $\mathcal{A}$, $at \subseteq \mathcal{A} \times \mathcal{A}$. Therefore, we say that $at(A, B)$ represents an attack from argument A to-

wards argument B . Let $\mathcal{L} : \mathcal{A} \rightarrow \{\text{in}, \text{out}, \text{undecided}\}$ be a labeling function (14). We define $at_{\mathcal{AF}}(A)$ as $\{B \mid at(B, A)\}$, where $A, B \in \mathcal{A}$. Let $in(\mathcal{L})$ be defined as $in(\mathcal{L}) = \{A \in \mathcal{A} \mid \mathcal{L}(A) = \text{in}\}$ and $out(\mathcal{L}), undecided(\mathcal{L})$ to be defined analogously. $\mathcal{L}$ is *admissible and complete* iff $\forall A \in in(\mathcal{L}), at_{\mathcal{AF}}(A) \subseteq out(\mathcal{L})$ and $\forall A \in out(\mathcal{L})$, there is B such that $B \in in(\mathcal{L})$ and $at(B, A)$ holds and $\forall A \in undecided(\mathcal{L}), \nexists B \in at_{\mathcal{AF}}(A), \mathcal{L}(B) = \text{in} \wedge \exists B' \in at_{\mathcal{AF}}(A), \mathcal{L}(B') \neq \text{out}$. Note that since $\mathcal{L}$ partitions $\mathcal{A}$ into exactly three sets, each direction of the standard bi-conditional follows from the other two conditions by exclusion. By adding external constraints we arrive at different semantics. Some common ones include Grounded Semantics, where $in(\mathcal{L})$ is minimal or Preferred Semantics where $in(\mathcal{L})$ is maximal.

3.2 ASPIC⁺ (4)

An ASPIC⁺ argumentation system is built over a logical language $\mathcal{L}$, closed under negation, together with a knowledge base $\mathcal{K} \subseteq \mathcal{L}$, partitioned into axiom premises $\mathcal{K}_n$, which are certain and not open to attack, and ordinary premises $\mathcal{K}_p$, which are defeasible. Inference is governed by a set of strict rules $\mathcal{R}_s$ of the form $\varphi_1, \dots, \varphi_n \rightarrow \psi$ and defeasible rules $\mathcal{R}_d$ of the form $\varphi_1, \dots, \varphi_n \Rightarrow \psi$.

An argument A over $\langle \mathcal{K}, \mathcal{R}_s \cup \mathcal{R}_d \rangle$ is constructed recursively: (i) if $\varphi \in \mathcal{K}$, then $A = \varphi$ is an argument with conclusion φ , premises $Prem(A) = \{\varphi\}$, and no subarguments; (ii) if $A_1, \dots, A_n$ are arguments with conclusions $\varphi_1, \dots, \varphi_n$ and $\varphi_1, \dots, \varphi_n \rightarrow \psi \in \mathcal{R}_s$ (resp. $\Rightarrow \psi \in \mathcal{R}_d$), then $A = A_1, \dots, A_n \rightarrow \psi$ (resp. $\Rightarrow \psi$) is an argument with conclusion ψ , premises $Prem(A) = \bigcup_i Prem(A_i)$, and immediate subarguments $Sub(A) = \{A_1, \dots, A_n\}$.

Arguments interact through three canonical attack types. An argument B *undermines* A if the conclusion of B negates an ordinary premise used in A ; B *rebuts* A if the conclusion of B negates the conclusion of a defeasible-rule application in A ; and B *undercuts* A if the conclusion of B negates the applicability of a defeasible rule used in A . Given the resulting set of arguments $\mathcal{A}$ and attack relation $at \subseteq \mathcal{A} \times \mathcal{A}$, the pair $\langle \mathcal{A}, at \rangle$ is precisely a Dung abstract argumentation framework in the sense of Section 3.1, and the labeling semantics established there apply to it directly.

3.3 Weighted Argumentation

To support quantitative reasoning over the structures introduced in Section 3.2, we adopt a weighted argumentation graph formalism, adapted from the weighted bipolar approach of (16) by restricting it to a single, weighted attack relation. A *WAF* is a tuple $\Gamma = \langle \mathcal{A}, R, \omega, \tau \rangle$, where $\mathcal{A}$ is a set of arguments, $R \subseteq \mathcal{A} \times \mathcal{A}$ is an attack relation, $\omega : R \rightarrow [0, 1]$ assigns each attack a weight expressing its severity, and $\tau : \mathcal{A} \rightarrow [0, 1]$ assigns each argument an initial strength. Unlike the qualitative labelling of Section 3.1, a WAF assigns every argument a graded strength value, with 0 denoting total rejection, 1 total acceptance, and intermediate values expressing degrees of acceptance. As in the bipolar case, we treat the computation of final strength as *modular*, following Mossakowski and Neuhaus (17): an aggregation function α combines the *weighted* strength of an argument's attackers – i.e. for each attacker B with $(B, A) \in R$, its contribution to the aggregate is a function of both $\tau(B)$ (or its current propagated strength) and $\omega(B, A)$ – and an influence function ι maps this aggregate, together with the argument's initial strength, to an updated strength. This separation of aggregation from influence is what allows the propagation mechanism described in Section 3.5 to be stated independently of how initial weights are obtained, and independently of how attack severities are assigned. In our case, initial strengths are drawn from the compositional intrinsic strength defined in Section 3.4, rather than from an atomic, externally supplied value, while attack weights ω remain exogenous, reflecting, e.g., the semantic force of the negation involved (undermining, rebutting, or undercutting) or a domain-specific severity judgment.

3.4 Intrinsic Strength

Standard WAF semantics leave the initial weight $\tau(a)$ of an argument unexplained, treating it as an exogenous input. Spaans (20) addresses this by studying principled methods for deriving an argument's *intrinsic strength* directly from the weights of the premises and inference rules used to construct it, rather than assigning it as an atomic value. We adopt this approach, using the premise weights $w(\varphi)$ and defeasible rule weights $w(r)$. The intrinsic strength $\sigma(A)$ of an argument A is defined recursively over its structure: if A is a premise argument, $\sigma(A) = w(\varphi)$ for its constituent premise φ ; if A is constructed by applying a rule r to subarguments $A_1, \dots, A_n$, then

$$\sigma(A) = \min(w(r), \sigma(A_1), \dots, \sigma(A_n)).$$

This weakest-link aggregation follows the possibilistic tradition (20) of treating an argument as no stronger than its least reliable component, and satisfies the monotonicity principles required by Spaans’s framework. Unlike approaches that derive initial strength from a model’s extraction confidence, $\sigma(A)$ is grounded exclusively in the rubric-assigned plausibility of A ’s premises and rules, avoiding the conflation between parsing reliability and evidential strength discussed in Section 2.

3.5 Gradual Semantics

To compute final acceptability degrees over arbitrary, potentially cyclic argument graphs, we require a propagation mechanism whose convergence does not depend on acyclicity. Discrete successive-update schemes, and continuous models such as the Quadratic Energy Model (18), guarantee convergence for acyclic graphs but offer, at best, empirical evidence of convergence on cyclic ones. We instead adopt the modular semantics framework of Mossakowski and Neuhaus (17), extended by Potyka (19), which derives convergence guarantees on arbitrary graphs, cyclic included, from the contraction mapping principle: if the influence function ι has a bounded derivative (Lipschitz constant strictly less than 1), the resulting update procedure is guaranteed to converge to a unique fixed point regardless of the graph’s cyclic structure. We select an influence function satisfying this bound, using $\sigma(A)$ (Section 3.4) as the initial weight in place of an atomic strength value. As noted by Potyka (19), convergence guarantees obtained this way trade off against *open-mindedness*, the extent to which final strength can move away from the initial weight, which we account for when selecting the Lipschitz bound.

4 Structural Overview

Figure 1 details PNsAV’s internal architecture, separating the pipeline into 4 distinct modules, interconnected by bidirectional data routes. The UI/UX layers acts as an intuitive interface for users, while also being the input/output module. The LLM extractor layer, represented by a ternary LLM orchestration, is tasked to parse the natural language into conventional ASPIC⁺ structure. Structural verification and graph data are executed by the symbolic validation layer, while the actual logic is computed by the engine, the final data being routed back to the user interface.




Figure 1: PNsAV's pipeline

5 Implementation

5.1 UI/UX

The PNsAV interface is fully integrated with the Python backend, using a multi-page setup built with a responsive two-column layout (`st.columns([1.5, 3.5])`) and a maximized (`layout="wide"`) screen view. The left panel handles data entry, allowing users to upload documents through (`st.file_uploader`) or type text directly into an (`st.text_area`). When you click the primary “Analyze” button (`st.button`), the app saves the data using (`st.session_state`) and starts processing it in the background. The right panel then shows the live results. The top section features an interactive graph built with `streamlit-agraph`, which renders arguments as colored nodes and connections as directed edges using a physics simulation. Directly below, the “Analyzed Text” section applies matching HTML highlights to ensure that key arguments share the exact same color codes as their corresponding nodes in the graph. At the same time, the “Logs” window updates automatically, using (`datetime.now().strftime`) to print the exact time and status dot next to every step the program takes. To give the app a professional look, we replaced the standard Streamlit design with custom CSS elements injected directly into the code. This allowed us to build a custom dark theme, rounded borders, specific box heights, and clean spacing optimized for long analysis sessions. Finally, web navigation is handled by a top bar split into columns: three buttons (“Resources”, “About”, and “Contact”) that let users switch pages by updating the session state, while the “PNsAV” logo on the left corner acts as a home button to instantly bring the user back to the main page.

5.2 Argument Mining

Natural language input is processed through a three-stage constrained LLM pipeline, structurally aligned with the ASPIC⁺ argument construction process introduced in Section 3.2: an *atom extractor* identifies premise-level propositions (annotated as axiomatic or normal premise, a *rule extractor* identifies inferential relations between them (annotated as strict or defeasible), and an *argument extractor* composes the resulting atoms and rules into structured arguments. Each stage is implemented as an independently prompted, guideline- and few-shot-constrained call, rather than a

single end-to-end extraction pass. This decomposition allows each stage’s output to be validated and constrained before being consumed by the next, isolating extraction errors to a specific stage rather than routing unconstrained free-text output directly into argument construction. Unlike approaches that fine-tune dedicated extraction models (e.g., (15)), this pipeline requires no labeled argument-mining dataset, at the cost of relying on prompt-level rather than parameter-level constraints on extraction quality. To account for this, we use chain-of-thought prompting (11).

5.3 The Validation Layer

PNsAV ensures that the extracted logic is not prone to crashing the logic layer by utilizing the Validation Layer. This layer unpacks the JSON provided by the LLM Extractor, which is confined in a strict schema, and verifies that the atoms, rules and arguments are constructed correctly (eg. an argument cannot have a conclusion different than the conclusion of its top rule). Using a basic DFS algorithm, the code checks for cycles, outputting a log to the user, displaying the graph type.

5.4 The Engine

5.5 Technical Aspects

6 Conclusions and future work

References and Notes

1. Olga Golovneva, Moya Chen, Spencer Poff, Martin Corredor, Luke Zettlemoyer, Maryam Fazel-Zarandi, and Asli Celikyilmaz. ROSCOE: A suite of metrics for scoring step by-step reasoning. *In Proceedings of the 11th International Conference on Learning Representations (ICLR)*, 2023.
2. Miller, Nate, "A Modular LLM Approach to Argument Extraction in Philosophical Texts" (2026). *Masters Theses*. 1170.

3. P. M. Dung. On the acceptability of arguments and its fundamental role in nonmonotonic reasoning, logic programming and n-person games. *Artificial Intelligence*, 77(2):321–357, 1995.
4. Henry Prakken. An abstract framework for argumentation with structured arguments. *Argument & Computation* 2010.
5. Jiayuan Mao, Chuang Gan, Pushmeet Kohli, Joshua B. Tenenbaum, and Jiajun Goel. The Neuro-Symbolic Concept Learner: Interpreting Scenes, Words, and Sentences from Natural Supervision. *In Proceedings of the 7th International Conference on Learning Representations (ICLR)*, 2019.
6. Nitish Gupta, Kevin Lin, Dan Roth, Mike Lewis, and Sameer Singh. Neural Module Networks for Reasoning over Text and Symbolic Data. *In Proceedings of the 8th International Conference on Learning Representations (ICLR)*, 2020.
7. Robin Manhaeve, Sebastijan Dumančić, Angelika Kimmig, Thomas Demeester, and Luc De Raedt. Neural probabilistic logic programming in DeepProbLog. *Artificial Intelligence*, 298:103504, 2021.
8. L. Cai, J. Choi, and J. He. Neural-Symbolic Commonsense Reasoner with Knowledge Graphs. *In Findings of the Association for Computational Linguistics (ACL)*, pages 2340–2352, 2021.
9. J. Tian, X. Zhang, and H. Wang. Weakly Supervised Neural Symbolic Learning for Cognitive Tasks. *In Proceedings of the 31st International Joint Conference on Artificial Intelligence (IJCAI)*, 2022.
10. Connor Pryor, Carlos Interian, and Luc de Raedt. Neuro-Symbolic Relational Reasoning Systems: A Survey. *arXiv preprint arXiv:2205.11904*, 2022.
11. Jundong Xu, Hao Fei, Liangming Pan, Qian Liu, Mong-Li Lee, Wynne Hsu Faithful Logical Reasoning via Symbolic Chain-of-Thought *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2024.

12. Liangming Pan Alon Albalak Xinyi Wang William Yang Wang LOGIC-LM: Empowering Large Language Models with Symbolic Solvers for Faithful Logical Reasoning *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023.
13. Anthony Hunter and Matthias Thimm. Probabilistic Argumentation with Epistemic Extensions and Incomplete Information. *arXiv:1405.3376 [cs.AI]*, 2014.
14. Caminada, M.W.A., Gabbay, D.M. A Logical Account of Formal Argumentation. *Stud Logica* 93, 109 (2009). <https://doi.org/10.1007/s11225-009-9218-x>
15. Gianvincenzo Alfano, Sergio Greco, Lucio La Cava, Stefano Francesco Monea, Irina Trubitsyna. LLM-based Argument Mining meets Argumentation and Description Logics: a Unified Framework for Reasoning about Debates. *arXiv:2603.02858 [cs.AI]*, march 2026.
16. Leila Amgoud and Jonathan Ben-Naim. Evaluation of arguments from support relations: Axioms and semantics. *Proceedings of the 26th International Joint Conference on Artificial Intelligence (IJCAI)*, 2017.
17. Till Mossakowski and Fabian Neuhaus. Modular Semantics and Characteristics for Bipolar Weighted Argumentation Graphs. *arXiv preprint arXiv:1807.06685*, 2018.
18. Nico Potyka. Continuous Dynamical Systems for Weighted Bipolar Argumentation. *Proceedings of the 16th International Conference on Principles of Knowledge Representation and Reasoning (KR)*, 2018.
19. Nico Potyka. Extending Modular Semantics for Bipolar Weighted Argumentation. *Proceedings of the 18th International Conference on Autonomous Agents and MultiAgent Systems (AAMAS)*, pages 1722–1730, 2019.
20. J. P. Spaans. Intrinsic Argument Strength in Structured Argumentation: a Principled Approach. *arXiv preprint arXiv:2109.00318*, 2021.